



Contact during exam:

Einar M. Rønquist (73593547 or 93404778)

**FINAL EXAM FOR COURSE TMA4280**  
**INTRODUCTION TO SUPERCOMPUTING**

Friday, May 31, 2010

Time: 09:00–13:00

Permitted aids: Approved calculator.  
All printed and hand written aids.

**In general:**

- Final grades will be announced by June 21.
- Short answers are fine as long as they are justified and explained.
- There are 15 questions and 4 points to be earned on each one.

**Problem 1**

Let  $\underline{A}$ ,  $\underline{B}$  and  $\underline{C}$  be real  $n \times n$  matrices. Consider the following operations:

$$\underline{C} = \underline{A} + \underline{B} \tag{1}$$

$$\underline{C} = \underline{A}\underline{B} \tag{2}$$

i.e., matrix addition and matrix multiplication, respectively. We are interested in the performance (in terms of the number of floating point operations per second) in the single processor case on `njord` (the current supercomputer at NTNU). The program is compiled using the highest level of optimization.

- a)** Will the performance of (1) and (2) be the same? If not, which operation will yield the highest performance? Explain your answer.

**Fasit:** Operation (2) will be faster than (1). The memory traffic for both operations is the same: we need to fetch  $\underline{A}$  and  $\underline{B}$  (each matrix has  $n^2$  floating point numbers), and we need to store  $\underline{C}$  ( $n^2$  floating point numbers). Operation (1) involves  $n^2$  flops (additions), while operation (2) involves approximately  $2n^3$  flops (multiplications and additions). Hence, the number of floating point operations per floating point number fetched or stored is  $\mathcal{O}(1)$  for operation (1) and  $\mathcal{O}(n)$  for operation (2). This yields a higher performance for operation (2), especially since we can also exploit the superscalar feature (chaining) available on each processor.

## Problem 2

We would like to solve the Poisson problem on the unit square  $\Omega = (0, 1) \times (0, 1)$  with boundary conditions  $u = 0$  specified along the boundary  $\partial\Omega$ . We discretize the problem using finite difference techniques; the Laplace operator is approximated using the standard 5-point stencil. The grid spacing in each spatial direction is  $h = 1/n$ , where  $n \gg 1$ .

The conjugate gradient method is used to solve the system of algebraic equations. The implementation is done in double precision. In the multi-processor case, the original grid is decomposed into  $P$  approximately equal subgrids, each of size  $\frac{n}{P} \times n$  (i.e., we "slice" the grid only in the  $x$ -direction), and we associate one subgrid with each processor.

In the following, you can assume that the time it takes to send a message of  $k$  bytes over the network can be expressed as  $\tau_C(k) = \tau_S + \gamma \cdot k$ , where  $\tau_S$  is the startup time and  $\gamma$  is the inverse bandwidth.

- a) Exactly how many unknowns do we have in this problem?

**Fasit:** The number of unknowns is precisely  $N = (n - 1)^2$  (the number of internal points).

- b) There are two contributions to the communication cost:  
 (i) the exchange of surface data between neighboring subgrids, and  
 (ii) the computation of global innerproducts.

For which value of  $P$  is the communication cost of (i) twice that of (ii)?

You can assume that  $n = 1000$ ,  $\tau_S = 10^{-6}$  s, and  $\gamma = 10^{-9}$  s/byte.

**Fasit:** From the lecture notes (and earlier exams), the communication cost per iteration for (i) is  $T_{comm,i} = 4\tau_C(8n)$ , while the communication cost per iteration for (ii) is  $T_{comm,ii} = 2\tau_S \log_2 P$ . For the first expression we have assumed that  $P > 2$ . Setting  $T_{comm,i} = 2T_{comm,ii}$  and using the given network model, we get

$$4(\tau_S + \gamma \cdot 8n) = 4\tau_S \log_2 P.$$

Inserting the given numbers yields  $\log_2 P = 9$  or  $P = 512$ . (Our earlier assumption that  $P > 2$  is thus valid.)

- c) Assume that for a particular value of  $n$ , and for  $P = 128$ , the total communication cost per conjugate gradient iteration is  $10^{-2}$  s. The parallel efficiency is equal to 0.5. The solution is obtained after 100 iterations. What is the total solution time (in seconds) on a single processor?

**Fasit:** From the lecture notes, the speedup can be expressed as

$$S_p = P \left( \frac{1}{1 + P \frac{T_{comm}^*}{T_1^*}} \right),$$

where  $T_{comm}^*$  is the communication cost per iteration and  $T_1^*$  is the solution time on a single processor per iteration. Since the parallel efficiency is equal to 0.5, we have  $P \frac{T_{comm}^*}{T_1^*} = 1$ . Hence,  $T_1^* = P T_{comm}^* = 128 \cdot 0.01s = 1.28s$ . With 100 iterations, the total solution time on a single processor is  $T_1 = 128s$ .

- d) How would you check whether the conjugate gradient iteration has converged?

**Fasit:** The residual vector  $\underline{r}$  is the difference between the right hand side and the left hand side of our system of equations. A small residual generally means that we have a good approximation of the solution. The length of the residual vector is automatically available during the conjugate gradient iteration (since we compute  $s = \underline{r}^T \underline{r}$ ). Hence, one alternative is to check if  $\sqrt{s} < tol$ , where  $tol$  is prescribed tolerance, e.g.,  $tol = 10^{-6}$ .

### Problem 3

For each point in this problem, please choose the correct alternative.

- a) Let  $a$ ,  $b$ , and  $c$  be three floating point numbers in double precision. Assume that  $a = 1 \cdot 10^{-12}$  and  $b = 2 \cdot 10^{-12}$ . We compute  $c = a + b$ . The answer  $c$  is accurate to about 4 digits.  
Answer: true or false.

**Fasit:** False. We have about 16 digits of accuracy.

- b) A recommended way to implement matrix-matrix multiplication on a single processor is to use an appropriate Level-1 BLAS routine.  
Answer: true or false.

**Fasit:** False. Level 1 BLAS include only vector operations. We should use Level 3 BLAS.

- c) It is faster to send one long message over the network rather than many small messages.  
Answer: true or false.

**Fasit:** True. This is because of the overhead associated with the startup of a message.

- d) The spectral bisection algorithm is useful in order to automatically partition a computational grid into subgrids. For two-dimensional problems, the subgrids will always cover approximately the same area.  
Answer: true or false.

**Fasit:** False. The spectral bisection algorithm will give us subgrids with approximately the same number of gridpoints, not necessarily the same area.

- e) In the multi-processor case, a global sum can be performed using `MPI_Allreduce` with the `MPI_SUM` option. The answer will be identical on all the processors, i.e., all the bits in the global sum will be the same on all the processors.  
Answer: true or false.

**Fasit:** True.

- f) Assume that  $P$  floating point numbers are evenly distributed across  $P$  processors. Using MPI, computing the product of all these numbers can be done equally fast as computing the sum.  
Answer: true or false.

**Fasit:** True.

- g) In the multi-processor case, the function `MPI_Scatter` is sometimes convenient to use. Do you expect to experience "deadlock" when using this function?  
Answer: yes or no.

**Fasit:** No. We expect the functions in the MPI library to be properly implemented and tested.

- h) It is typically easier to achieve good speedup for a code compiled with the highest level of compiler optimization compared to a code using no compiler optimization at all.  
Answer: true or false.

**Fasit:** False. See lecture notes, exercises and earlier exams.

- i) On `njord` (the current supercomputer at NTNU), it is recommended to exclusively use OpenMP if you need to run your program on more than 32 processors.

Answer: true or false.

**Fasit:** False. We can only use shared memory programming within a single node on `njord`, i.e., using a maximum of 16 processors. We need to use MPI when we need to communicate across nodes.

- j) The first part of a parallel program involves reading input data from a single file stored on a single hard drive. In the single-processor case, the program spends about 1% of the time on this task. The maximum speedup for this program is 100.

Answer: true or false.

**Fasit:** True. Reading from a single file is a sequential operation, even in the multiple processor case. Hence, the sequential part of this program will at least be 1%. According to Amdahl's law, the maximum speedup is 100.

Good luck!

Einar M. Rønquist